

Toxicity Classification in Civil Comments

Problem:

Toxicity classification in Civil comments/ online conversations on a public platform, focused on news articles.

A NLP task where comment texts are processed to learn the toxicity in them, to classify whether a comment is toxic or not; where toxicity is defined as anything rude, disrespectful or otherwise likely to make someone leave a discussion.

As part of a [competition](#) on Kaggle, the Conversation AI team, a research initiative founded by Jigsaw and Google, released the data labeled for toxicity. Crowdsourcing was used for annotation.

The problem can be divided into two parts: building a model to learn to classify toxicity in the comments; and mitigating the unintended model bias for identity terms(e.g. muslim, gay. Background: [here](#))

For this project, we focus primarily on building a deep stacked Bi-LSTM model to train for the toxicity classification.

Data:

[1.8 million civil comments labeled for toxicity](#)

Initialization Details:

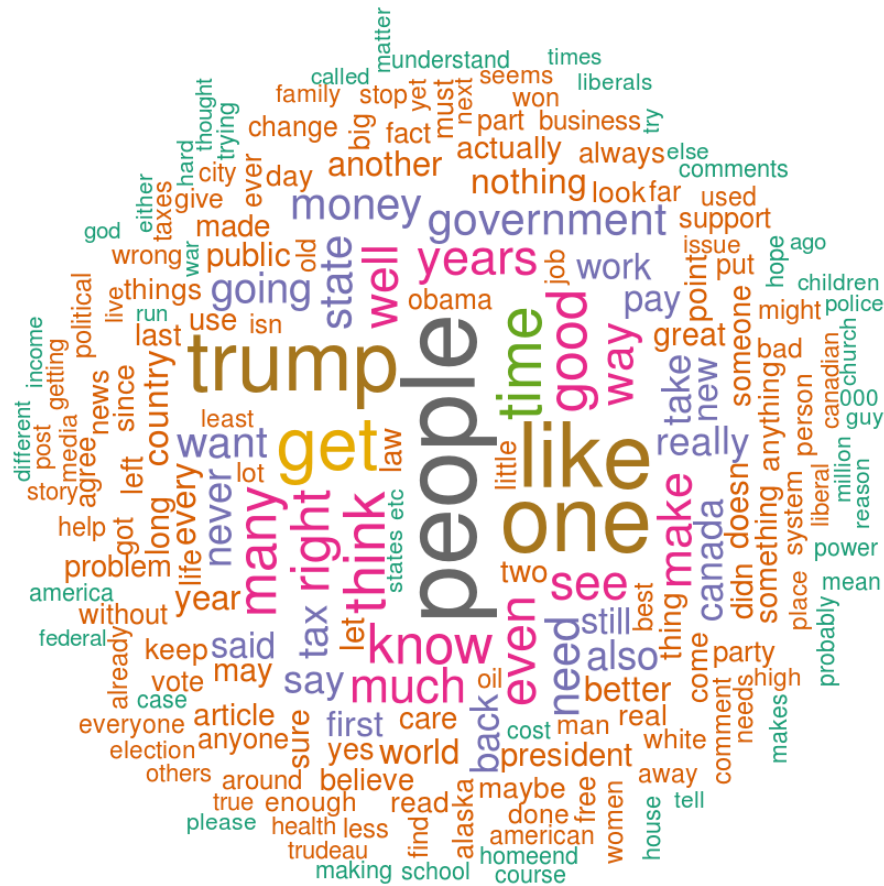
0.75/0.25 train and test division used for this project.

300d GloVe word embeddings pretrained on Wikipedia and Gigaword are used for input feature representation.

Dropout of 0.2 used on each LSTM layer. Adam optimizer is used.

Random Uniform initialization for Embedding weights. Glorot Uniform initialization for all hidden kernel weights, and Orthogonal initialization for recurrent kernel weights in LSTM units. Binary Cross Entropy loss function is used, as this is a binary classification task.

Data Insights:



Pipeline used: Clean special chars | tokenize | remove stop words | filter words less than size 2

Parameters used for cloud: min_freq:1, max_words: 200

Word frequencies:

Word	Count
people	206209
one	178437
like	178381
trump	160343
get	137280
time	109246

Data sample:

comment_text	target
This is so cool. It's like, 'would you want your mother to read this??' Really great idea, well done!	0
Thank you!! This would make my life a lot less anxiety-inducing. Keep it up, and don't let anyone get in your way!	0
This is such an urgent design problem; kudos to you for taking it on. Very impressive!	0
Is this something I'll be able to install on my site? When will you be releasing it?	0
haha you guys are a bunch of losers.	0.893617
ur a sh*tty comment.	0.666667
hahahahahahahahhha suck it.	0.457627

Model Architecture:

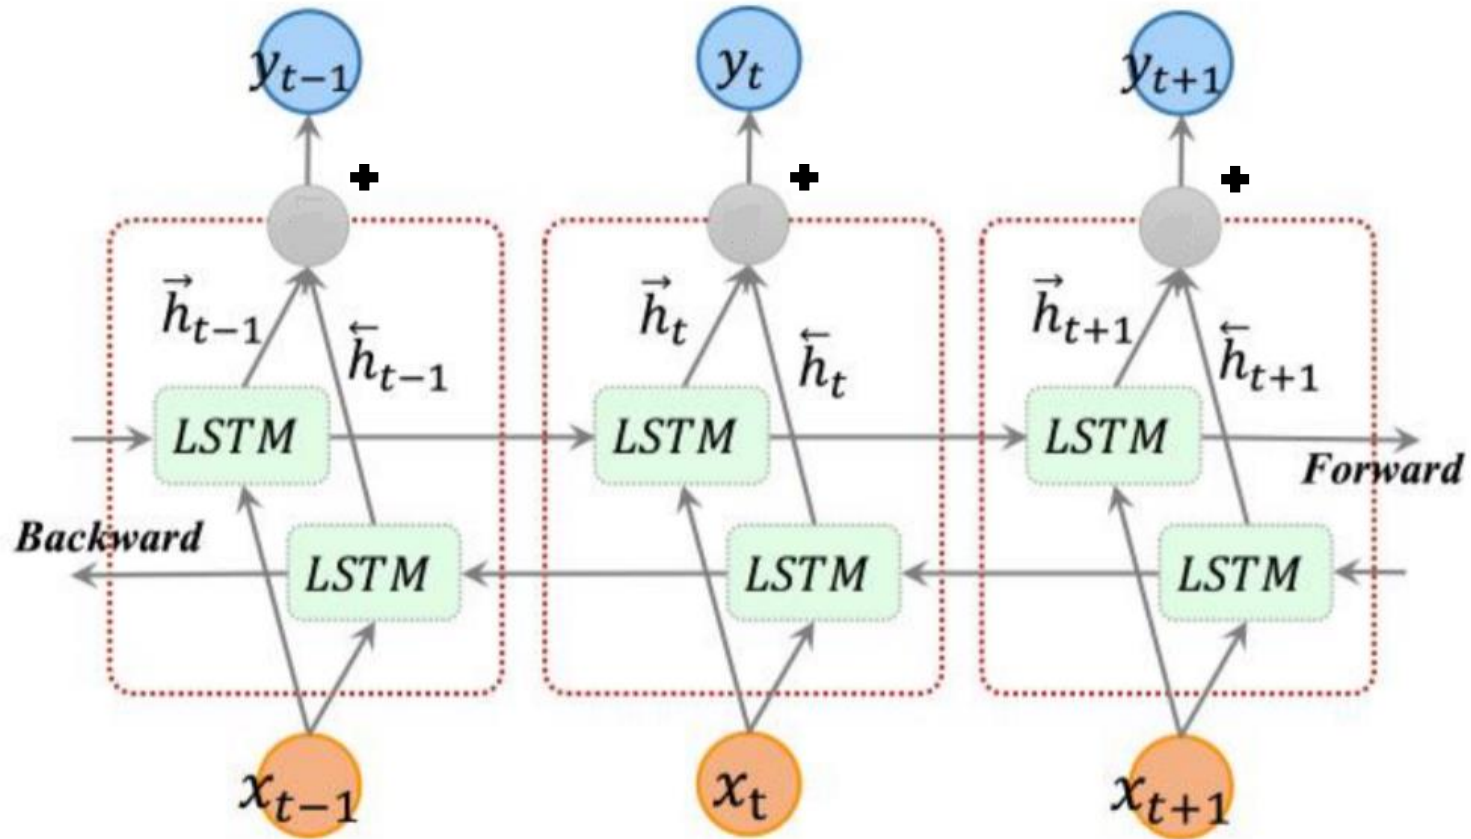


Figure 1a. BiLSTM internal architecture ([courtesy](#))




Figure 1b. Stacked BiLSTM

Model:

Stacked bidirectional LSTM model, with 2 BiLSTM layers, a ReLu activated fully connected layer on top, and an sigmoidal output layer. Figure 1b. gives an overview of our model.

We used Keras CuDNNLSTM layer while constructing the model, which implements a standard LSTM layer. Fig 1a. gives an overview of the BiLSTM structure. The input gate, the forget gate, the output gate and the input cell state, which are represented in the LSTM cell in Fig. 1a, can be represented as:

$$\begin{aligned} f_t &= \sigma_g(W_f x_t + U_f h_{t-1} + b_f) \\ i_t &= \sigma_g(W_i x_t + U_i h_{t-1} + b_i) \\ o_t &= \sigma_g(W_o x_t + U_o h_{t-1} + b_o) \\ \tilde{C}_t &= \tanh(W_C x_t + U_C h_{t-1} + b_C) \end{aligned}$$

where Wf , Wi , Wo , and WC are the weight matrices mapping the hidden layer input to the three gates and the input cell state; while the Uf , Ui , Uo , and UC are the weight matrices connecting the previous cell output state to the three gates and the input cell state. The bf , bi , bo , and bC are four bias vectors. The σg is the gate activation function, and the \tanh is the hyperbolic tangent function. From these, the cell output state and layer output state can be represented as:

$$\begin{aligned} C_t &= f_t * C_{t-1} + i_t * \tilde{C}_t \\ h_t &= o_t * \tanh(C_t) \end{aligned}$$

Finally, in the BiLSTM layer concatenates these operations are performed in both directions, forward and backward on the sequences, thus generating two sets of hidden outputs., which are concatenated to be passed as inputs to the next layer. Each hidden output can be represented as:

$$y_t = \text{concatenate}(h_t^{\rightarrow}, h_t^{\leftarrow})$$

Minimal tuning done on the hyperparameters, and results shown are using the optimal parameters found during the process.

Model summary

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 150, 300)	92703600
bidirectional_7 (Bidirection	(None, 150, 200)	321600
dropout_7 (Dropout)	(None, 150, 200)	0
bidirectional_8 (Bidirection	(None, 200)	241600
dropout_8 (Dropout)	(None, 200)	0
dense_7 (Dense)	(None, 64)	12864
dense_8 (Dense)	(None, 1)	65
Total params: 93,279,729		
Trainable params: 576,129		
Non-trainable params: 92,703,600		

Training Results

Train on 1353655 samples, validate on 451219 samples
Epoch 1/4
1353655/1353655 [=====] - 490s 362us/step - loss: 0.1480 - acc: 0.9450 - val_loss: 0.1303 - val_acc: 0.9497

Epoch 2/4
1353655/1353655 [=====] - 494s 365us/step - loss: 0.1250 - acc: 0.9513 - val_loss: 0.1261 - val_acc: 0.9503

Epoch 3/4
1353655/1353655 [=====] - 494s 365us/step - loss: 0.1187 - acc: 0.9532 - val_loss: 0.1229 - val_acc: 0.9517

Epoch 4/4
1353655/1353655 [=====] - 494s 365us/step - loss: 0.1136 - acc: 0.9546 - val_loss: 0.1232 - val_acc: 0.9522